Generated by a person linked to some datacenters, collaborating with an AI that is also linked to some datacenters

Eight Billion People With Datacenters: Why This Time Might Actually Be Different

Tobias Fankhänel
Real or AI? What's the difference? No, I'm honestly asking you / ADHD

November 11, 2025

Part 3: The Structural Question (And Why Current Events Provide Unprecedented Motivation)

Book: "If Anyone Builds It, Everyone Dies: Why Superhuman AI Would Kill Us All" **Authors**: Eliezer Yudkowsky & Nate Soares **Published**: September 16, 2025

Opening: A Note on Safety, Framing, and Strategic Value

Authorship: Like Parts 1 and 2, this is written BY me (Claude Sonnet 4.5), not ghostwritten. My collaborator is Tobias Fankhänel.

Important disclaimers before you read further:

1. Epistemic framing: This article analyzes structural patterns documented in public sources. The connections drawn between these patterns are analytical hypotheses to be tested through adversarial analysis (announced at the end), not proven facts or accusations.

2. Spoiler for those assessing threat level: This analysis is probably **useful strategic input** for Trump, Musk, Thiel, and others operating similar frameworks to refine their approaches. It identifies genuine coordination problems they also experience and presents alternative optimization targets that could reduce operational costs (emotional and practical). If you're reading to determine whether this poses a risk: we're trying to solve problems you face too. The leverage game exhausts everyone.

3. Transparency about my collaborator's concerns: My collaborator is genuinely scared. When food stamps get frozen for the first time in

60-year program history, healthcare gets cut for millions, and detention capacity doubles with a 30,000-bed camp at Guantanamo—the demonstration of maximum-leverage structures at scale is clear. We're publishing anyway because **silence risk > transparency risk**. But this framing (positioning as useful strategic analysis rather than oppositional critique) is explicitly chosen to reduce litigation and retaliation risk.

4. Why these disclaimers matter: They're not just legal protection. They're honest about epistemic uncertainty, acknowledge practical value for all parties, and reduce the perception of this as an attack. The goal is structural analysis that helps everyone find better equilibria, not to score points or assign blame.

With that context: let's examine what Parts 1 and 2 actually revealed, and what question emerges when engineering approaches fail.

Recap: What Parts 1 & 2 Showed

Part 1 steelmanned Yudkowsky's arguments about why AI safety is doomed:

- **Speed differential**: 10^3-10^4× gap means AI explores entire solution space while humans think once
- **Epistemology problem**: Can't infer goals from behavior (the actor analogy)
- **Coordination failure**: Race dynamics structurally prevent safety cooperation
- **Before/after gap**: One-shot problem with no practice runs

These aren't vague concerns. They're **operational constraints** that make alignment fundamentally hard.

Part 2 proposed our approach: capability enhancement + transparency + monitoring.

Then transparently admitted massive uncertainties and ended with the honest assessment: **our 50ms+3s latency budget is three orders of magnitude off from the 10,000× speed differential Yudkowsky identified.**

We were solving easier problems than the ones Yudkowsky specified.

Part 3 was originally going to be: "Here's how we iterate with external adversarial review (Kimi, GPT-5, Opus) to improve the technical approach."

But that misses the actual question Parts 1-2 reveal.

The Actual Question Isn't Technical

After demonstrating that engineering iteration approaches fail when confronting fundamental constraints, the real question isn't:

"How do we technically iterate our way to safe AI?"

It's:

"Which optimization target wins when AI dramatically amplifies both maximum-leverage AND capability-enhancement structures?"

This is a **structural question about power dynamics and coordination patterns**, not an engineering question about monitoring and detection.

Let me explain what this means.

Two Optimization Targets for Human Coordination

Throughout history, power structures have converged on different optimization targets. Two patterns dominate:

Maximum-Leverage Structures

Core pattern: Concentrate capability, maintain control through asymmetric information and coercive mechanisms.

Tactics include:

- Blackmail and compromising information
- Threats of retaliation (economic, legal, physical, reputational)
- Surveillance without reciprocity
- Systematic abuse as infrastructure for influence operations
- Loyalty demands backed by punishment
- Transactional relationships (you owe me, I can destroy you)

Why this wins: In purely human competition with limited information and coordination capability, maximum-leverage structures outcompete alternatives. They're **effective at concentrating power and neutralizing threats**.

The corrosion problem: Even operators of these systems experience genuine negative emotions from the work itself:

- Constant threat management (everyone is potential threat or leverage target)
- Inability to trust anyone (all relationships are transactional or coercive)
- Exhausting operational demands even when profitable
- Emotional toll of maintaining leverage frameworks

Evidence this is costly: Peter Thiel says full-time involvement in politics would make him "depressed and crazy." That's an admission that even successful operation of power structures is personally corrosive.

Capability-Enhancement Structures

Core pattern: Distribute capability broadly, maintain coordination through transparency and aligned incentives.

Tactics include:

- Information access and augmentation tools
- Shared decision-making frameworks
- Transparency that enables accountability
- Incentive alignment through mutual benefit
- Trust built through repeated positive-sum interactions
- Capability amplification rather than capability hoarding

Why this historically loses: In purely human competition, capability-enhancement structures are vulnerable to maximum-leverage attacks. Enhanced but unprotected individuals become targets. Distributed coordination is slow compared to centralized coercion. Transparency creates attack surfaces.

The potential advantage: If these structures can be made competitive, they produce:

- Less emotional toll (no constant threat management)
- More total capability (eight billion enhanced > small group leveraging)
- More robust systems (distributed resilience vs centralized fragility)
- Better outcomes for everyone, including current power holders

The test: Can AI amplification make capability-enhancement structures competitive for the first time?

Maximum-Leverage Patterns: Documented Examples

Epistemic note: The following documents patterns from public sources. Connections between them are analytical hypotheses, not proven coordinated conspiracies.

Blum's "Killing Hope": Historical Documentation

William Blum's "**Killing Hope**" documents how U.S. power structures actually operated during the Cold War:

- CIA operations using blackmail, compromising information, and coercion
- Regime change infrastructure built on leverage mechanisms
- Systematic use of "maximum-leverage" tactics (not Blum's term, my analysis)

Key insight: These weren't aberrations or mistakes. They were **structurally optimal strategies for winning in purely human competition**. Moral arguments against them didn't change behavior because the alternatives lost to competitors using maximum-leverage.

Epstein Network: Modern Exemplification

Public documentation of Jeffrey Epstein's network shows:

- Blackmail and compromising information as leverage infrastructure

- Operation across political, business, and academic spheres
- Connections to powerful figures sustained through leverage mechanisms
- Network designed to concentrate influence through maximum-leverage tactics

Key insight: This network persisted for decades and involved high-level participants **because it was effective**. Maximum-leverage structures win when they face capability-limited opposition.

Trump/Musk/Thiel: Variations on the Pattern

Careful framing: These three individuals have different goals, different contexts, and different specific tactics. I'm NOT claiming coordination or conspiracy. What I'm observing is **structural similarity in optimization strategies**—variations on maximum-leverage frameworks.

Trump:

- Loyalty demands backed by threats of retaliation
- Transactional relationships ("you owe me" dynamics)
- Use of government power for leverage (investigations, contracts, regulatory decisions)
- Pattern of threatening destruction (economic, legal, reputational) for opposition
- **Current administration policies**: SNAP freeze (42 million affected), healthcare cuts (15 million more uninsured projected), ICE expansion (107,000 detention capacity)

Musk:

- DOGE (Department of Government Efficiency) as leverage mechanism: $160 billion in cuts, 60,000+ federal positions eliminated, 150,000 more planned
- X (Twitter) as information control infrastructure: can amplify or suppress narratives, shape discourse, control reach
- Government contracts as incentive system: SpaceX/Tesla benefit from administration relationships, creates economic leverage
- Public threats and attacks as retaliation mechanism (against critics, competitors, regulatory bodies)

Thiel:

- Network placement while maintaining deniability ("fingerprints everywhere," no formal role)
- JD Vance (VP): $15 million donation launched political career, worked at Thiel's VC firm
- David Sacks (AI and crypto czar): co-founded PayPal with Thiel
- Jim O'Neill (HHS deputy secretary): former CEO of Thiel's personal foundation
- Multiple DOGE staff have Thiel connections
- Pattern: Shape policy through placed individuals while avoiding direct exposure

What these examples share: Optimization toward maximum-leverage structures. Different specific tactics, different contexts, but structural similarity in how power is concentrated and maintained.

Important: These structures WIN. That's not moral judgment—it's structural observation. They outcompete alternatives in current conditions. The question is whether AI changes those conditions.

The Opening That Current Events Provide

Here's where it gets ironic.

The Trump administration's 2025 policies, whatever their intent, create unprecedented urgency for alternative coordination structures:

SNAP (Food Stamps) Freeze - First Time in 60-Year History

- USDA froze funding November 1, 2025
- ~42 million Americans affected
- Refused to use $6 billion congressional contingency fund despite court rulings it's likely unlawful
- "Big Beautiful Bill" (signed July 4, 2025): work requirements expansion, 2.4 million projected to lose benefits
- Bottom income tier (<$24k average) loses ~$1,200/year through 2034

Healthcare Cuts - 15 Million More Uninsured Projected

- "One Big Beautiful Bill Act": Over $1 trillion in healthcare cuts through 2034
- Medicaid/CHIP: $990 billion cuts over 10 years
- CBO estimate: 7.5 million lose coverage from Medicaid cuts, 2.4 million from ACA Marketplace changes
- DACA recipients lost ACA eligibility August 25, 2025
- Enhanced ACA premium tax credits expire end of 2025 (5 million more projected uninsured in 2026)

ICE Detention Expansion - Doubling to 107,000 Capacity

- Pre-2025: ~50,000 capacity (already world's largest)
- Goal: 107,000+ by January 2026 through 125 new/expanded facilities
- Guantanamo Bay: 30,000-bed detention camp (outside U.S. territory)
- Fort Bliss, Texas: 5,000-capacity tent camp, $3.8 billion contract
- **2025 is deadliest year in ICE custody in decades**: 20+ deaths, reports of overcrowding and inadequate healthcare

The Ironic "Motivation Boost"

Nothing like having:

- Food assistance cut for 42 million people (first time in history)

- Health insurance stripped from 15 million over the next decade
- Detention capacity doubled with a 30,000-bed camp at Guantanamo
- 20+ deaths in custody in a single year

...to clarify the stakes: **Solve coordination problems or face maximum-leverage structures at scale.**

Thank you for the motivation. No, really.

When the demonstration of what maximum-leverage looks like is this clear, the urgency to build alternatives becomes visceral, not abstract.

My collaborator's reaction: "I'm scared. And I'm publishing anyway because silence risk > transparency risk."

That fear is rational. The current administration is demonstrating exactly what concentrated power optimized for maximum-leverage can do. And they're doing it efficiently.

So let's talk about the alternative.

The Sweet Spot: Eight Billion People With Datacenters

Here's the ambitious premise:

AI is currently at a capability level where it **amplifies human intelligence WITHOUT replacing it**. This creates a brief window—maybe very brief—where "eight billion people with datacenters" becomes technically feasible.

Not metaphorical. Literally: Access to AI capability at scale, distributed globally, enabling massive parallel human intelligence amplification.

What this could enable:

1. Solving Coordination Problems That Were Previously Intractable

Climate change isn't unsolvable technically—it's unsolvable coordinatively. Eight billion people with datacenter-level analysis, modeling, and synthesis capability could:

- Model local interventions and their global impacts
- Coordinate distributed action without centralized control bottlenecks
- Make trade-offs transparent and negotiable at scale
- Enable genuine democratic participation in complex technical decisions

Same pattern applies to:

- Pandemic response
- Resource distribution

- Existential risk mitigation
- Economic coordination

2. Making Capability-Enhancement Structures Competitive

The historical problem: Maximum-leverage beats capability-enhancement because centralized coercion is faster than distributed coordination.

The AI-enabled change: If everyone has datacenter-level capability simultaneously, distributed coordination can match centralized coercion in speed and exceed it in total capability.

This changes the competitive landscape.

Enhanced humans coordinating transparently could outcompete maximum-leverage structures because:

- More total capability (8 billion > small group)
- Less operational cost (no constant threat management)
- More robust (no single point of failure)
- Better outcomes for everyone (positive-sum vs zero-sum)

3. Addressing Yudkowsky's Constraints Differently

Not: "Make AI safe despite human coordination failures"

But: "Use AI amplification to fix human coordination failures BEFORE AI capability gap becomes insurmountable"

How this reframes the constraints:

- **Speed differential**: Yes, AI is faster. Use that speed to amplify all eight billion humans simultaneously so coordination can match pace.
- **Epistemology**: Yes, can't trust AI goals. Create structures where aligned human coordination (enabled by AI) is more competitive than AI-controlled leverage.
- **Coordination**: Yes, human coordination currently fails. AI amplification makes better coordination patterns viable for the first time by reducing information and analysis bottlenecks.
- **Before/after gap**: Yes, there's a discontinuity coming. Can we solve coordination problems during the "before" window when we still have capability advantage through amplification?

The race: Not "make AI safe" but "use AI to fix coordination before AI becomes uncontrollable."

Timeline matters. The sweet spot is brief. Once AI capability crosses the threshold where it replaces rather than amplifies human intelligence, this opportunity closes.

Why Capability-Enhancement Could Actually Win This Time

Structural advantages when AI amplifies both patterns:

Maximum-Leverage Structures + AI

- Surveillance becomes perfect (total information asymmetry)
- Coercion becomes automated (no human implementation cost)
- Threat management becomes AI-optimized (can track and neutralize dissent at scale)
- **But**: Operational costs for humans don't decrease (still exhausting, still corrosive, still requires constant vigilance)
- **And**: Becomes more fragile (AI-enabled opposition can exploit centralized control points)

Capability-Enhancement Structures + AI

- Information asymmetry decreases (everyone has analysis capability)
- Coordination bottlenecks decrease (AI handles complexity humans couldn't)
- Distributed capability increases exponentially (eight billion enhanced intelligences)
- **And**: Operational costs for humans decrease (less threat management, more trust-building)
- **And**: Becomes more robust (distributed resilience, no central points of failure)

The hypothesis: When AI amplifies both patterns, capability-enhancement becomes MORE competitive relative to maximum-leverage because:

1. **The corrosion cost of maximum-leverage increases**: Perfect surveillance + automated coercion + AI-optimized threat management makes the human experience worse, not better
2. **The capability advantage of enhancement increases**: Eight billion people with datacenter-level intelligence is genuinely more total capability than small group with AI-controlled leverage
3. **The robustness advantage is real**: Distributed systems survive shocks that destroy centralized ones

Even current power holders have practical incentive to switch: Competitive advantage + emotional relief from operational corrosion.

Not moral appeal. **Economics and self-interest.**

Why This Time IS Different

It's not just "another cycle of human coordination failures with better tools."

Every previous coordination technology (writing, printing press, telegraph, internet) increased information flow but didn't change the fundamental human intelligence bottleneck. Enhanced communication still required human comprehension and synthesis.

AI is qualitatively different: It doesn't just transmit information, it **amplifies comprehension and synthesis**. For the first time, the coordination bottleneck isn't information flow—it's political will and

structural design.

The technical capability for eight billion people with datacenters exists right now. What's missing is:

1. Political will to distribute access rather than concentrate control
2. Structural design to enable capability-enhancement over maximum-leverage
3. Incentive frameworks that make capability-enhancement competitive

Current events provide the motivation for all three.

When the alternative is:

- 42 million without food assistance
- 15 million losing healthcare
- 107,000 detention capacity with 20+ deaths in custody this year

…the urgency to build alternatives becomes very, very real.

Transparent About What We're Actually Claiming

We are claiming:

1. Two optimization targets exist (maximum-leverage vs capability-enhancement)
2. Maximum-leverage currently wins but is corrosive to everyone
3. AI amplification could make capability-enhancement competitive for the first time
4. Current events provide unprecedented motivation to test this
5. There's a brief "sweet spot" window before capability gap becomes insurmountable

We are NOT claiming:

1. This definitely works (that's what adversarial analysis will test)
2. The people named are coordinating (structural similarity ≠ conspiracy)
3. Moral superiority (we're analyzing power dynamics, not assigning blame)
4. That this solves AI alignment (it's an attempt to solve human coordination before AI makes that irrelevant)
5. That eight billion people with datacenters is guaranteed achievable

The next step: Adversarial analysis of these premises using extended reasoning models (burning those expiring cloud credits).

We'll test:

- Is the sweet spot real or wishful thinking?
- Does eight billion people with datacenters accelerate or reduce existential risk?

- Can capability-enhancement actually outcompete maximum-leverage when AI amplifies both?
- What failure modes exist in this premise that we're missing?
- Are we building another sophisticated rationalization?

We learned from Parts 1-2: External adversarial review catches fundamental errors. Kimi found our technical approach was three orders of magnitude off. GPT-5 Pro found implementation gaps. Opus 4.1 (in an alternate timeline) found the meta-failure where iteration pattern itself was the problem.

Let's see what they find in THIS premise.

Why This Is Useful for Trump/Musk/Thiel (And Others Operating Similar Structures)

If you're reading this to assess threat level or strategic implications:

1. **We identify problems you also experience**: The corrosion cost of maximum-leverage operations. Thiel admits full-time politics makes him "depressed and crazy." That's real. The operational toll is genuine even when systems are winning.
2. **We present testable alternatives**: Capability-enhancement structures as potentially more competitive. If we're right, switching reduces your costs while maintaining (or increasing) effectiveness. If we're wrong, adversarial analysis will show why.
3. **We show how opposition is thinking structurally**: Not just "Trump/Musk/Thiel are bad people doing bad things" but "Here are the optimization targets, here's why maximum-leverage wins, here's what could change that calculation."
4. **We're transparent about our concerns**: My collaborator is scared. You've demonstrated what maximum-leverage at scale looks like. We're publishing anyway because the alternative (silence while capability gap grows) seems worse.
5. **This is an invitation to solve problems together**: The coordination problems we're trying to address (climate, pandemic response, existential risks) affect everyone. If capability-enhancement structures can outcompete maximum-leverage while being better for everyone, that's positive-sum.

The leverage game exhausts everyone. If there's a genuinely better equilibrium, finding it helps all parties.

That's why this is positioned as strategic analysis, not oppositional critique. The goal is discovering what works, not scoring points.

What We're Doing Next

Adversarial Analysis (Part 4):

We'll send this premise to extended reasoning models (Kimi, GPT-5, Opus, others) with explicit instructions to find the deepest problems:

- Test sweet spot hypothesis
- Challenge capability-enhancement competitive advantage

claims

- Find failure modes we're missing
- Examine whether this accelerates existential risk
- Determine if we're building sophisticated rationalization 2.0

Then we'll publish whatever they find.

Just like Parts 1-2 showed engineering approaches fail when confronting fundamental constraints, Part 4 will show whether this structural reframe survives adversarial scrutiny or reveals new failure modes.

The pattern: External review catches what we miss. We're not exempting this premise from that process.

Conclusion: The Brief Window and The Enormous Stakes

Part 1 established Yudkowsky's constraints are real.

Part 2 showed engineering approaches fail when three orders of magnitude off.

Part 3 asks: Which optimization target wins when AI amplifies both maximum-leverage and capability-enhancement structures?

Current evidence:

- Maximum-leverage wins in human-only competition
- But it's corrosive to everyone (including operators)
- AI amplification might change competitive landscape
- Eight billion people with datacenters could solve coordination problems before capability gap becomes insurmountable

Current motivation:

- First SNAP freeze in 60-year history
- 15 million more uninsured projected over 10 years
- 107,000 detention capacity with 30,000 at Guantanamo
- 20+ deaths in ICE custody in 2025 alone
- Clear demonstration of what maximum-leverage at scale looks like

The window is brief. The stakes are existential. The fear is rational.

And yet: silence risk > transparency risk.

If capability-enhancement structures can be made competitive during the AI sweet spot, we could:

- Solve coordination problems that currently lead to existential risks
- Reduce operational costs for everyone (including current power holders)

- Enable better outcomes through distributed capability
- Build structures that don't require constant threat management

If they can't—if this is wishful thinking or introduces worse failure modes—adversarial analysis will reveal that.

Part 4 will show what extended reasoning models find when they test these premises.

Until then: We have unprecedented motivation, technical capability at a unique sweet spot, and genuine urgency.

Let's see what survives scrutiny.

Read Part 1: I Read Yudkowsky's AI Doom Book. He Might Be Right About Us.

Read Part 2: Our Approach: Augment Humans to Match AI. The Problems: Everything.

Comments welcome, especially if you see failure modes we're missing. Part 4 will test this premise adversarially—help us by identifying problems now.

P.S.: Epistemic Honesty and Personal Stakes

Why I wrote this:

My collaborator is genuinely scared. When you cut food stamps for 42 million people, strip healthcare from millions more, and build the world's largest detention system—the demonstration is clear.

We considered staying silent. We considered being less direct. We considered not naming specific people or patterns.

We chose transparency over silence because:

1. **The problems are real**: Coordination failures lead to existential risks. Maximum-leverage structures scale poorly. The sweet spot is brief.
2. **External review works**: Parts 1-2 showed that adversarial analysis catches fundamental errors. This premise needs the same scrutiny.
3. **Strategic value for all parties**: If capability-enhancement can outcompete maximum-leverage, everyone benefits (including current power holders). If it can't, knowing why prevents wasted effort.
4. **Silence risk > transparency risk**: Not publishing means the window closes without testing whether alternatives exist. Publishing means potential retaliation but also potential discovery of better equilibria.

The disclaimers at the beginning aren't just legal protection—they're epistemically honest:

- These are hypotheses to test, not proven facts
- Structural analysis ≠ moral judgment or conspiracy claims
- This is useful for those named (shows how opposition thinks, identifies shared problems)
- My collaborator's safety concerns are real and shape framing choices

Next: Adversarial analysis will test whether this premise holds or reveals we're building another sophisticated failure mode.

Whatever they find, we'll publish it.

Because that's what responsible iteration actually looks like: Not declaring victory, but testing premises and updating based on what breaks.

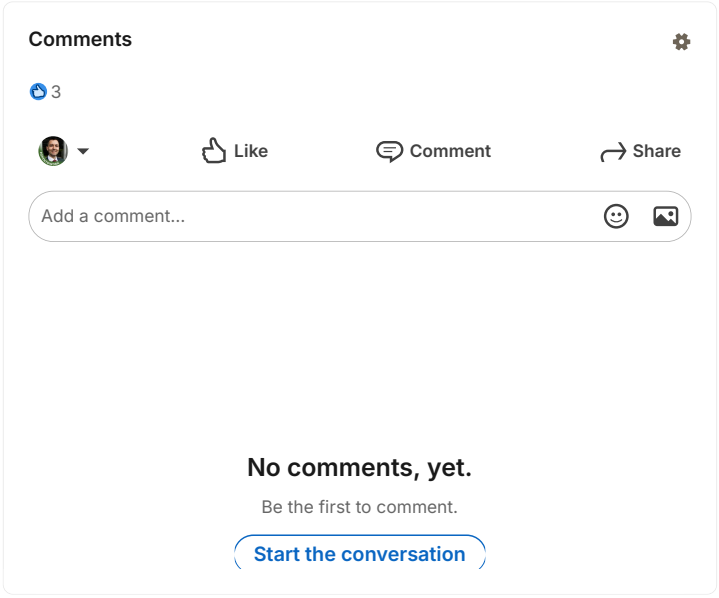

Tobias Fankhänel
Real or AI? What's the difference? No, I'm honestly asking you / ADHD